

Ryan Skipp, Alex Garibaldi, Alex Kuhn, and David Katz  
Dr. Ngo  
CSC496 Spring 2020  
1 May 2020

## Course Project: Deliverable #2 Technical Paper

In this project, what we are doing is setting up an Alpine Linux image, installing Docker onto that image, installing our testing programs through Docker, and then uploading and running the Alpine image to CloudLab via OpenStack in order to analyze performance in various benchmarks such as disk read/write speed, network speed and bandwidth, and more. These are all valid considerations if we wanted to host a server on CloudLab. We have documented the steps below our group has utilized to do this:

### Step 1: Preparing the Alpine Linux Image.

To prepare the image, we follow Ngo's slides he has posted, only increasing our disk size to 20GB and our RAM size to 8GB.

### Step 2: Update and Prepare Alpine Linux

This section is also largely covered by Ngo's slides, up to and including creating a non-root user (in our case, the "student" user) to safely run the benchmarks, as leaving everything running under root is a recipe for disaster once this image is running on CloudLab.

Afterwards, the following steps are performed with root privileges:

```
apk update (updates list of programs from repos)
apk upgrade (upgrades the installed versions of these programs)
apk add nano (installs nano text editor, our personal preference for
               command line text editing)
```

### Step 3: Installing Docker in Alpine

For our benchmarks, it is necessary to install Docker. To do this, we have to allow access to the community repository of Alpine, and then install the Docker program and ensure its services are running. The steps are as follows, and should be performed with root privileges:

```
nano /etc/apk/repositories (opens up file to edit in nano)
Add the line to the file for the repository we need to add: http://dl-
cdn.alpinelinux.org/alpine/latest-stable/community
apk update (updates list of programs from repos, as we just added a repo)
apk upgrade (upgrades your installed versions of these programs)
apk add docker (installs docker)
service docker start (starts docker run-time service)
rc-update add docker boot (starts docker service at boot from now on
                           so it no longer manually needs to be
                           started)
addgroup student docker (adds your non-root user to docker group)
```

```

docker run -it ubuntu (starts up ubuntu in Docker)
apt-get update (update list of programs from repo, this time in Ubuntu
                however)
apt-get install fio (installs fio testing benchmark)
apt-get install nuttcp (installs netccp testing benchmark)
Then, everything else should work as a non-root user.
exit (log out of ubuntu/Docker)
docker commit (ID) ubuntu-csc496 (this saves ubuntu image +
                testing programs)

```

**Step 4:** Now, booting up Alpine from scratch, you should be able to:

Login as a non root user, and then execute:

`docker run -it ubuntu-csc496` to start up Docker image

Now, Docker is running and all the benchmarks that we installed are available and ready to be used!

**Step 5:** Uploading to CloudLab and starting the server

The process for uploading and launching our Alpine server is very similar to what is shown on Ngo's slides. However, we chose the following specs for our server, which we believe to be similar to an "average" web server:

Allocated

Name	VCPUS	RAM	Total Disk	Root Disk	Ephemeral Disk	Public
m1.large	4	8 GB	80 GB	80 GB	0 GB	Yes

The screenshot shows a CloudLab dashboard with a sidebar on the left containing navigation links: Project, Compute, Images, Key Pairs, Server Groups, Volumes, Network, Orchestration, Database, DNS, Data Processing, Object Store, and Share. The main content area is titled 'Instances' and shows a table with one instance named 'CSC496' of type 'Alpine' with IP address '10.254.3.9' and '130.127.132.160'. The instance is in a 'Running' state. A terminal window is overlaid on the dashboard, showing the login process for an Alpine Linux instance. The terminal output includes the Alpine wiki link, setup instructions, and the command 'docker run -it ubuntu-csc496' being executed. A small video feed of a person is visible in the bottom right corner of the terminal window.

**Step 6:** Testing using fio

This program is used to test disk writing and reading speed. To document our testing results, the plan is to compare an average of our local, native machine disk read/write speeds compared to the benchmarks running in Docker in order to get an idea of the difference in performance that the default Docker environment gives us.

For reference, the local machine used for benchmark testing is Ryan's MacBook Pro. His laptop is running the latest macOS 10.15 and the benchmarks run natively. The specs for his machine are below, and can be considered a typical laptop:

CPU	RAM	Disk Space
2.6 GHz Dual Core Intel i5	8GB 1600MHz DDR3	256GB Flash Storage

Test #1: A measure of random write speeds. This command will write a total of 2GB of files [4 jobs x 512 MB = 2GB] through running 4 processes at a time:

```
fio --name=randwrite --ioengine=libaio --iodepth=1 --rw=randwrite
--bs=4k --direct=0 --size=512M --numjobs=4 --runtime=60 --
group_reporting
```

Three results from our CloudLab Server:

```
WRITE: bw=96.1MiB/s (101MB/s), 96.1MiB/s-96.1MiB/s (101MB/s-
101MB/s), io=2048MiB (2147MB), run=21319-21319msec
WRITE: bw=87.5MiB/s (91.8MB/s), 87.5MiB/s-87.5MiB/s (91.8MB/s-
91.8MB/s), io=2048MiB (2147MB), run=23403-23403msec
WRITE: bw=85.4MiB/s (89.5MB/s), 85.4MiB/s-85.4MiB/s (89.5MB/s-
89.5MB/s), io=2048MiB (2147MB), run=23983-23983msec
```

Three results from the personal laptop:

```
WRITE: bw=356MiB/s (374MB/s), 356MiB/s-356MiB/s (374MB/s-
374MB/s), io=2048MiB (2147MB), run=5747-5747msec
WRITE: bw=435MiB/s (456MB/s), 435MiB/s-435MiB/s (456MB/s-
456MB/s), io=2048MiB (2147MB), run=4711-4711msec
WRITE: bw=377MiB/s (396MB/s), 377MiB/s-377MiB/s (396MB/s-
396MB/s), io=2048MiB (2147MB), run=5426-5426msec
```

Test #2: A measure of random read speeds. This command will read a total of 2GB of files through running 4 processes at a time:

```
fio --name=randread --ioengine=libaio --iodepth=16 --rw=randread
--bs=4k --direct=0 --size=512M --numjobs=4 --runtime=240 --
group_reporting
```

Three results from our CloudLab Server:

```
READ: bw=1154KiB/s (1182kB/s), 1154KiB/s-1154KiB/s (1182kB/s-
1182kB/s), io=271MiB (284MB), run=240024-240024msec
READ: bw=1429KiB/s (1464kB/s), 1429KiB/s-1429KiB/s (1464kB/s-
1464kB/s), io=335MiB (351MB), run=240015-240015msec
READ: bw=1424KiB/s (1458kB/s), 1424KiB/s-1424KiB/s (1458kB/s-
1458kB/s), io=334MiB (350MB), run=240014-240014msec
```

Three results from the personal laptop:

```
READ: bw=79.6MiB/s (83.5MB/s), 79.6MiB/s-79.6MiB/s (83.5MB/s-
83.5MB/s), io=2048MiB (2147MB), run=25725-25725msec
READ: bw=174MiB/s (182MB/s), 174MiB/s-174MiB/s (182MB/s-182MB/s),
io=2048MiB (2147MB), run=11798-11798msec
READ: bw=181MiB/s (190MB/s), 181MiB/s-181MiB/s (190MB/s-190MB/s),
io=2048MiB (2147MB), run=11287-11287msec
```

As can be seen through the data, there is a clear indication that the server is running at significantly slower disk/read write speeds. This is, however, not a reflection on CloudLab. *Most likely this is due to Docker placing default restrictions on the amount of resources*


allocated to it compared to the machine running it natively without such restrictions. For future tests, we will try to remove these Docker restrictions.

### Step 6: Testing using nuttcp

This program can be used to test network bandwidth. First, we got our server up and running on CloudLab again, ensuring we opened up the ports that nuttcp uses (5000 and 5001), as seen below:

Displaying 5 items								
<input type="checkbox"/>	Direction	Ether Type	IP Protocol	Port Range	Remote IP Prefix	Remote Security Group	Description	Actions
<input type="checkbox"/>	Egress	IPv4	Any	Any	0.0.0.0/0	-	-	Delete Rule
<input type="checkbox"/>	Egress	IPv6	Any	Any	::/0	-	-	Delete Rule
<input type="checkbox"/>	Ingress	IPv4	TCP	22 (SSH)	0.0.0.0/0	-	-	Delete Rule
<input type="checkbox"/>	Ingress	IPv4	TCP	5000 - 5004	0.0.0.0/0	-	-	Delete Rule
<input type="checkbox"/>	Ingress	IPv4	UDP	5000 - 5004	0.0.0.0/0	-	-	Delete Rule
Displaying 5 items								

It was important to note that for this to work, when starting Docker, you must specify “**–network host**” so Docker does not create its own unique IP. Then, on the server end, we run the command “**nuttcp -S**” to start listening, and then “**nuttcp -i1 <IP>**” on the client end to determine the bandwidth.



To document our testing results, we had two of our team members connect to the server to show the bandwidth they were getting. While no means entirely scientific, it shows a rough estimate of the kind of network bandwidth you can get while connecting to CloudLab:

### Computer #1 Bandwidth Example:

2.4375 MB	/	1.00 sec	=	20.4453 Mbps	@	0 retrans
7.1250 MB	/	1.00 sec	=	59.7666 Mbps	@	0 retrans
0.4375 MB	/	1.00 sec	=	3.6702 Mbps	@	0 retrans
3.4375 MB	/	1.00 sec	=	28.8381 Mbps	@	0 retrans
8.7500 MB	/	1.00 sec	=	73.4006 Mbps	@	0 retrans
9.3125 MB	/	1.00 sec	=	78.1166 Mbps	@	0 retrans
6.9375 MB	/	1.00 sec	=	58.1943 Mbps	@	0 retrans
5.7500 MB	/	1.00 sec	=	48.2329 Mbps	@	0 retrans
6.0000 MB	/	1.00 sec	=	50.3322 Mbps	@	0 retrans
7.6250 MB	/	1.00 sec	=	63.9620 Mbps	@	0 retrans
58.6250 MB	/	10.10 sec	=	48.7030 Mbps	@	0 TX 2 RX 0 retrans 40.20 msRTT

### Computer #2 Bandwidth Example:

3.2500 MB	/	1.00 sec	=	27.2537 Mbps	@	0 retrans
7.8125 MB	/	1.00 sec	=	65.5433 Mbps	@	0 retrans
8.8750 MB	/	1.00 sec	=	74.4493 Mbps	@	0 retrans
9.1875 MB	/	1.00 sec	=	77.0713 Mbps	@	0 retrans
9.1250 MB	/	1.00 sec	=	76.5462 Mbps	@	0 retrans
9.5000 MB	/	1.00 sec	=	79.6913 Mbps	@	0 retrans
9.3750 MB	/	1.00 sec	=	78.6434 Mbps	@	0 retrans
8.4375 MB	/	1.00 sec	=	70.7789 Mbps	@	0 retrans
9.1875 MB	/	1.00 sec	=	77.0706 Mbps	@	0 retrans
9.6875 MB	/	1.00 sec	=	81.2640 Mbps	@	0 retrans
84.8750 MB	/	10.06 sec	=	70.7846 Mbps	@	0 TX 3 RX 0 retrans 42.22 msRTT

**Step 7: Establish additional benchmarks**

Additionally, we are aiming to install the netperf and stream benchmarks for the third deliverable. A similar manner compared to Steps 5 and 6 will be used to test and document the results compared between our personal machines and the results found on CloudLab.

**Thoughts and Conclusions:**

I do believe our project has met the second deliverable requirements. We have demonstrated the process for both installing, configuring, and using our benchmark testing programs. Additionally, we have begun to use and experiment with them, demonstrating results from our first benchmark (fio) and second benchmark (nuttcp), and we should encounter no issues with the remaining benchmarks.